

CSCE 633: Machine Learning

Lecture 13: Logistic Regression and Gradient Descent


Texas A&M University

Bobak Mortazavi


Goals

- Re-motivate Logistic Regression
- Understand why the shift from regression to classification results in needing Gradient Descent
- Learn how to use Cross-Validation

Classification – Step Functions




one-dimensional input:
decision boundary is a single point




two-dimensional input:
decision boundary is a line

Linear Regression for Classification



If you set $y > 0.5$ as a decision rule – not clear it works
Hard to interpret – results are not contained in $[0,1]$

Step function from Linear Regression Boundary



Bernouli Distribution

- We can create a probability density function that represents a single experiment asking yes/no

$$Y \sim \text{Bernouli}(\theta), Y \in \{0,1\}$$

$$p(y|\theta) = \theta^{I(y=1)}(1 - \theta)^{I(y=0)} = \begin{cases} \theta, & y = 1 \\ 1 - \theta, & y = 0 \end{cases}$$

- Can think of this as a coin toss experiment and the likelihood of heads vs. tails

Logistic Regression

- **Parametric classification** method (**not regression**), which is sometimes referred to as a “generalization” of linear regression because
 - We still compute a linear combination of feature inputs, $\mathbf{x}^T \mathbf{w}$ (sometimes written $\mathbf{w}^T \mathbf{x}$)
 - However, instead of estimating the continuous output variable, we pass this into a function

$$\mu(\mathbf{w}^T \mathbf{x}) = \frac{1}{1 + e^{-(\mathbf{w}^T \mathbf{x})}}$$

- Where $0 \leq \mu(\mathbf{w}^T \mathbf{x}) \leq 1$, and where the Gaussian noise of linear regression is replaced by the Bernoulli Distribution so that

$$p(y|\mathbf{x}, \mathbf{w}) = \text{Ber}(y|\mu(\mathbf{w}^T \mathbf{x}))$$

- Therefore, the output belongs to a class 1 ($y = 1$) with probability $\mu(\mathbf{w}^T \mathbf{x})$, and class 0 ($y = 0$) with probability $1 - \mu(\mathbf{w}^T \mathbf{x})$

Why use a Sigmoid Function?


$$\sigma(\eta) = \frac{1}{1 + e^{-(\eta)}} = \frac{e^{\eta}}{1 + e^{\eta}}$$

- Has very nice properties for classification
 - Bounded between 0 and 1 <- thus interpretable as a probability
 - Monotonically increasing <- thus can be used for classification rules

$\sigma(\eta) > 0.5$, positive class ($y = 1$)


$\sigma(\eta) \leq 0.5$, negative class ($y = 0$)

- Nice computational properties for optimizing criterion function



Decision Boundaries (Margin Classification)

Types of Loss



Logistic Regression: Representation

- Setup classification problem for two classes
 - Input: $\mathbf{x} \in \mathbb{R}^p$
 - Output: $y \in \{0, 1\}$
 - Training Data: $D = \{(x_1, y_1), \dots, (x_n, y_n)\}$
 - Model parameters: \mathbf{w} (weights)
 - Model:

$$p(y | \mathbf{x}, \mathbf{w}) = \sigma(\mathbf{w}^T \mathbf{x}), \sigma(\eta) = \frac{1}{1 + e^{-(\eta)}} = \frac{e^\eta}{1 + e^\eta}$$

$$y = \begin{cases} 1, & p(y | \mathbf{x}, \mathbf{w}) > 0.5 \\ 0, & \text{otherwise} \end{cases}$$

Maximum Likelihood Estimation

$$\hat{\theta} = \underset{\theta}{\operatorname{argmax}} L(\theta; x)$$

- Likelihood is $\prod_{p:y_p=1} p(x_p) \prod_{p':y_{p'}=0} (1 - p(x_{p'}))$
- With one dimension this means optimizing intercept w_0 and slope w_1
- So maximizing the parameters means finding the optimal w that maximize the product

Multiple Logistic Regression

$$\log\left(\frac{p(x)}{1 - p(x)}\right) = w_0 + w_1x_1 + \cdots + w_px_p$$

- $p(y|x, w) = \text{Ber}(y|\sigma(w^T x))$, with a linear decision boundary at $p(x) > 0.5$
- We can then calculate the negative of the log of the likelihood (negative log likelihood)

$$NLL(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

Multiple Logistic Regression

$$\log\left(\frac{p(x)}{1 - p(x)}\right) = w_0 + w_1x_1 + \cdots + w_px_p$$

- $p(y|x, w) = \text{Ber}(y|\sigma(w^T x))$, with a linear decision boundary at $p(x) > 0.5$
- We can then calculate the negative of the log of the likelihood (negative log likelihood)

$$NLL(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

- Now suppose we recast $\tilde{y} = \{-1, 1\}$, then

$$NLL(w) = - \sum_{i=1}^n \left(\log \left(1 + e^{-\tilde{y}_i w^T x_i} \right) \right)$$

- Which has no closed form solution

Logistic Regression: Evaluation

- Data likelihood (1 training sample)

$$p(y|x) = \begin{cases} \sigma(w^T x) & y = 1 \\ 1 - \sigma(w^T x) & \text{otherwise} \end{cases} = \sigma(w^T x)^y (1 - \sigma(w^T x))^{1-y}$$

- Data likelihood (all training samples)

$$L(D, w) = \prod_{i=1}^n p(y_i|x_i, w) = \prod_{i=1}^n \sigma(w^T x_i)^{y_i} (1 - \sigma(w^T x_i))^{1-y_i}$$


- Log Likelihood

$$l(D, w) = \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

- Negative Log Likelihood (the cross-entropy error)

$$nll(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

Training Logistic Regression



Where training in linear regression shifted line (slope/intercept)
training in logistic regression alters the “step”-like nature of the sigmoid

Training Logistic Regression with Gradient Descent

- Gradient Descent Methods
- How to optimize the coefficients

Logistic Regression: Optimization

- The cross-entropy error

$$\mathcal{E}(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

- Must find weights w of logistic regression that minimize the error
- Because there is no closed form solution, unlike linear regression, we need gradient descent

$$\nabla \mathcal{E}(w) = \sum_{i=1}^n (\sigma(w^T x_i) - y_i) x_i$$

- Is this convex?

Cross Entropy Error: Convex

$$\nabla \mathcal{E}(w) = \frac{\partial \mathcal{E}(w)}{\partial w} = \sum_{i=1}^n (\sigma(w^T x_i) - y_i) x_i$$

$$H = \frac{\partial^2 \mathcal{E}(w)}{\partial w \partial w^T} = \sum_{i=1}^n \sigma(w^T x_i) (1 - \sigma(w^T x_i)) (x_i x_i^T)$$

For all \mathbf{v} , substituting $\mu = \sigma(w^T x_i) (1 - \sigma(w^T x_i)) \geq 0$

$$\mathbf{v}^T H \mathbf{v} = \mu \mathbf{v}^T \left(\sum_{i=1}^n x_i x_i^T \right) \mathbf{v} = \mu \sum_{i=1}^n (x_i^T \mathbf{v})^T (x_i^T \mathbf{v}) = \mu \sum_{i=1}^n \|x_i^T \mathbf{v}\|_2^2 \geq 0$$

Gradient Descent in Logistic Regression

- The cross-entropy error

$$\mathcal{E}(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

- Gradient Descent Optimization Expression

$$w \leftarrow w - \alpha(k) \nabla \mathcal{E}(w)$$
$$\nabla \mathcal{E}(w) = \sum_{i=1}^n (\sigma(w^T x_i) - y_i) x_i$$

Batch Gradient Descent

1. Initialize $w, \epsilon, \alpha, k = 0$
2. While $\|\nabla \mathcal{E}(w)\|_2 > \epsilon$
 - 2a. $k = k + 1$
 - 2b. $w = w - \alpha(k)(\sum_{i=1}^n (\sigma(w^T x_i) - y_i) x_i)$

Stochastic Gradient Descent

Update weights one sample at a time

1. Initialize $w, \epsilon, \alpha, k = 0$
2. Loop until convergence
 - 2a. $k = k + 1$
 - 2b. Randomly choose a single sample (x_i, y_i)
 - 2c. Compute its contribution to the gradient $g_i = (w^T x_i - y_i)x_i$
 - 2d. Update weights $w = w - \alpha(k)g_i$

Mini Batch Gradient Descent

Update weights with a subset of samples at a time

1. Initialize $w, \epsilon, \alpha, k = 0$
2. Loop until convergence
 - 2a. $k = k + 1$
 - 2b. Randomly choose a set of samples (x_i, y_i)
 - 2c. Compute their contribution to the gradient $g_i = (w^T x_i - y_i)x_i$
 - 2d. Update weights $w = w - \alpha(k) \sum_{i=1}^M g_i$


This is a good compromise between batch gradient descent (complete gradient) and stochastic gradient descent (gradient of one sample at a time)

Common mini-batch is a set M between 50 and 250 samples

This is typically used for training neural networks

What if data does not fit a line?

Example: Samples of a sine function



We can use a non-linear basis function

$$\phi(x): x \in \mathbb{R}^p \rightarrow z \in \mathbb{R}^M$$


So that $y = w^T z = w^T \phi(x)$

More non-linear basis functions


Example: Samples from a sine function

Polynomial basis function $\phi(\mathbf{x}) = [1 \ x \ \dots \ x^M]^T$


$M = 0$ underfitting




$M = 1$ underfitting



$M = 3$



$M = 9$ overfitting



But higher-order non-linearity tends to overfit


Weights of high order polynomials are very large

$$y_i = \mathbf{w}^T \mathbf{z}_i = \mathbf{w}^T \phi(\mathbf{x}_i), \quad \mathbf{z}_i = \phi(\mathbf{x}_i) \in \mathbb{R}^M$$

	$M = 0$	$M = 1$	$M = 3$	$M = 9$
w_0	0.19	0.82	0.31	0.35
w_1		-1.27	7.99	232.37
w_2			-25.43	-5321.83
w_3			17.37	48568.31
w_4				-231639.30
w_5				640042.26
w_6				-1061800.52
w_7				1042400.18
w_8				-557682.99
w_9				125201.43

Overfitting

- The risk of using highly flexible (complicated) models without enough data is that the model overfits
- Leads to **poor generalization**
- How do you detect overfitting?
 - Can plot model complexity vs. the objective function
 - As complexity improves, will see training improves while testing improves then deteriorates



- How do you prevent this?
 - Feature selection, more data, or regularization

Linear Regression: Summarization

- Representation

$$y = w^T x$$

- Evaluation

$$\min_w RSS(w) = \sum_{i=1}^N (y_i - w^T x_i)^2$$

- Analytic Optimization

$$w^* = (X^T X)^{-1} X^T y$$

- Approximate Optimization

via gradient descent

- In linear regression the maximum likelihood estimation = the ordinary least squares solution

Logistic Regression: Summarization

- Generalization of linear regression

$$\sigma(w^T x)$$

is interpretable as a probability

- Evaluation through cross-entropy error

$$\mathcal{E}(w) = - \sum_{i=1}^n (y_i \log \sigma(w^T x_i) + (1 - y_i) \log(1 - \sigma(w^T x_i)))$$

- Optimization through gradient descent

Practical Use of Logistic Regression

- Understanding re-sampling
- Measures of Accuracy when repeating modeling
- Python/Discussion!

Resampling Methods. Why?

- Resampling is a method by which we re-draw subsets of training data to see if model fit changes
- This allows us to see how stable our model is by how much the coefficient weights change
- We want to reduce model error rates
- How does one guarantee performance on a test set when only evaluating a training set?

Resampling Methods

Cross-validation

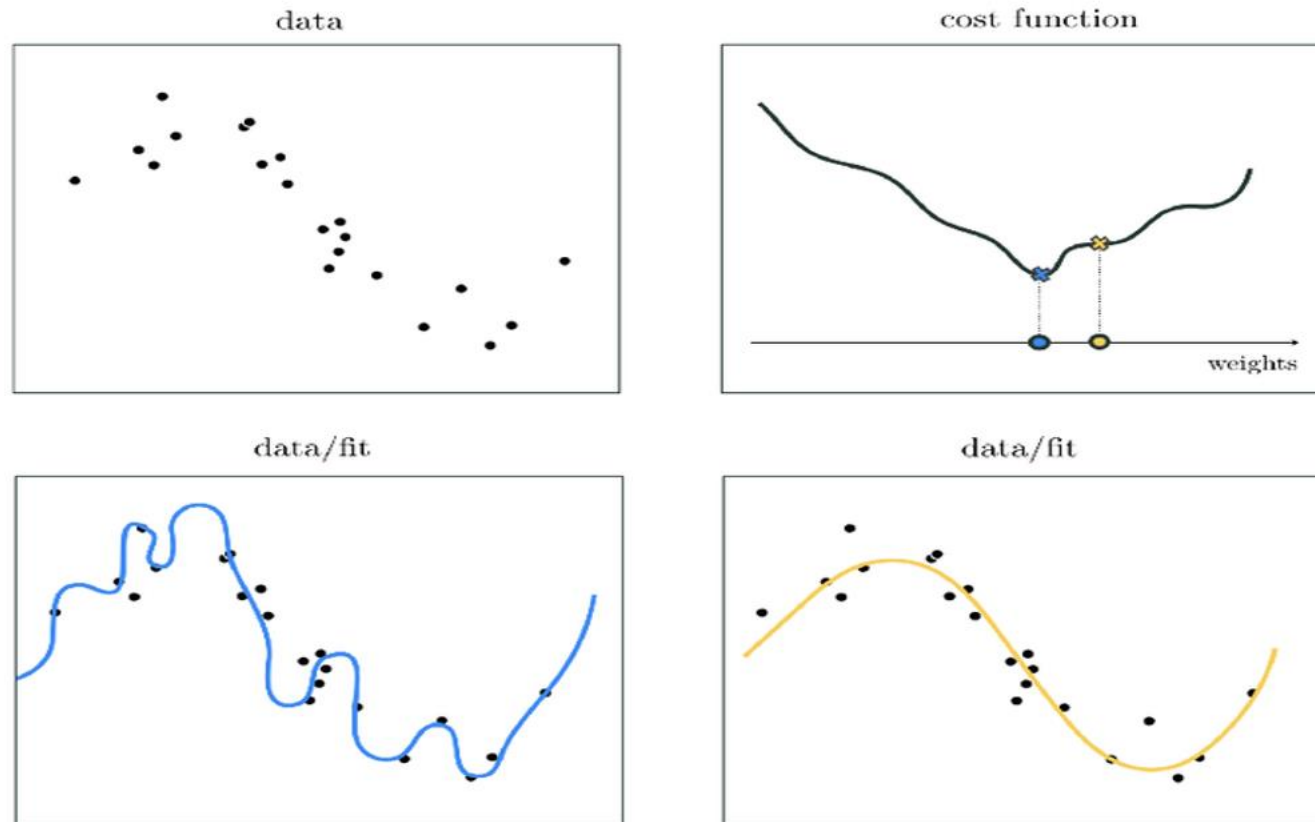
- Helps estimate measures of performance
- Allows for a choice in levels of model flexibility
- Used for hyperparameter tuning
- (You tend to see this more in computer science work)

Bootstrapping

- Also helps give estimates of model performance
- Measures of accuracy surrounding parameter estimates
- Repeated trials help simulate different event rates
- (you tend to see this more in biostatistics work)


Overfitting

- Consider fitting a polynomial model to data, how does training error change with increasing degree?



Overfitting

- As a model gets more complex, we become prone to reducing training error at the expense of testing error
- We also see sensitivity increase to specific inclusion of variables




Cross-Validation: Test Set Approach

- We can randomly split the training data (after having removed our “hold out” data) to train and test (or validate).
- These splits can be 50/50, 80/20, 90/10




Let's compare polynomial fit for horsepower and mpg

- Model for estimating relationship between engine horsepower and the miles per gallon achieved
- Best fit is $mpg = w_0 + w_1 \text{horsepower} + w_2 \text{horsepower}^2 + \epsilon$
- So how do we determine the best degree of polynomial?




Cross-validation: Test set approach

- Best fit is $mpg = w_0 + w_1horsepower + w_2horsepower^2 + \epsilon$
- Split the data 50/50
- See that the cubic term actually increases error




Cross-validation: Test set approach

- Best fit is $mpg = w_0 + w_1horsepower + w_2horsepower^2 + \epsilon$
- Split the data 50/50
- See that the cubic term actually increases error
- What if we repeat this experiment multiple times?




Cross-validation: Test set approach

- Best fit is $mpg = w_0 + w_1horsepower + w_2horsepower^2 + \epsilon$
- Split the data 50/50
- See that the cubic term actually increases error
- What if we repeat this experiment multiple times?



Leave-one-out Cross-validation




Advantages of LosoCV


- In each iteration, train with as much data as possible
- Consider error as the average across all iterations
- This presents far less bias
- It won't overestimate/underestimate test error as much as a single train/test split can

Comparison with MPG and Horsepower

Train-test split




Loso cv




K-fold Cross-Validation

- LOSO CV can be computationally expensive. A single train/test split not robust enough.
- Iterate subsets instead to find a balance




Comparison with MPG and Horsepower

Train-test split



K=10 fold cross-validation



Takeaways

- Why do we need Logistic Regression?
- Understand why the shift from regression to classification results in needing Gradient Descent
- Learn how to use Cross-Validation
- **Next time: Regularization**